

JVM 调优经验分享：让你的 Java 应用性能更上一层楼

/ 如何用一行代码修复一个排查了一周的问题

主讲人：朱孟柱

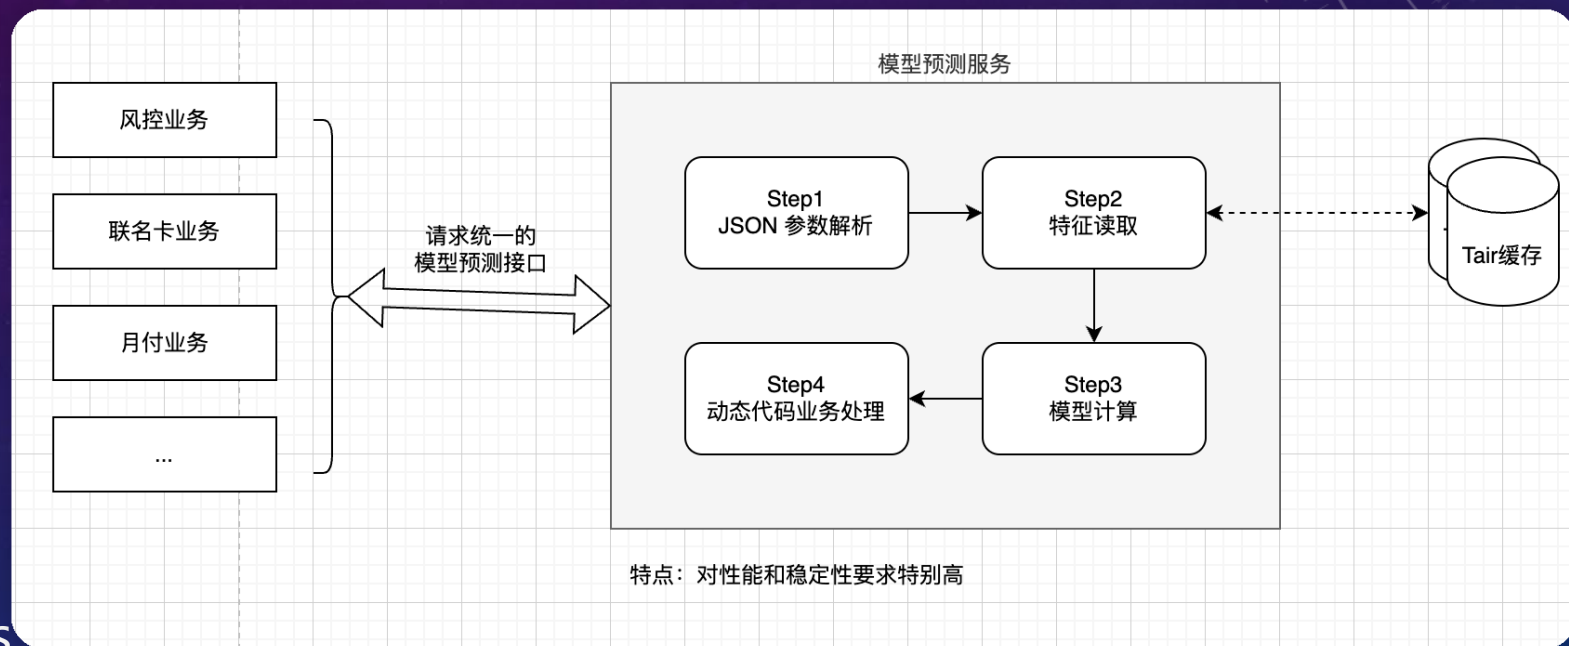
目标

通过这个案例，了解：

- 排查性能问题的一些基本思路；
- G1GC 的基本原理；
- Eclipse Memory Analyzer 的基本使用方法；
- JVM 内存 dump 方法及分析思路；

问题：服务请求成功率从 99.99% 下降到 99%

- 模型预测服务
- 语言：Java8, G1GC
- 机器：4C8G * 40
- QPS：峰值 1W
- SLA：99.99%
- 性能：TP9999 \leq 50ms



问题：服务请求成功率从 99.99% 下降到 99%

- 模型预测服务
- 语言：Java8, G1GC
- 机器：4C8G * 40
- QPS：峰值 1W
- SLA：99.99%
- 性能：TP9999 \leq 50ms

```
1 public class ModelPredictService {
2
3     /**
4      * 伪代码
5      */
6     public <T> T doExecute(String sceneId, String paramJsonStr) {
7         // 1. 解析 JSON 参数
8         final Map<String, Object> paramMap = JacksonJsonParser.parse(paramJsonStr);
9         // 2. 调用特征查询接口
10        final Map<String, Double> featureMap = FeatureQuery.query(sceneId, paramMap);
11        // 3. 调用模型预测接口
12        final double score = ModelPredict.predict(featureMap);
13        // 4. 调用结果处理接口
14        return ResultProcess.process(score, featureMap, paramMap);
15    }
16
17 }
```


问题：服务请求成功率从 99.99% 下降到 99%

首轮排查：

- 最近没有代码变更
- 最近没有重启服务
- 最近没有修改机器配置
- 排查了服务器网络状况，正常
- 排查了 Tair 缓存服务状况，正常
- 排查了服务器和 Tair 之间的网络拓扑，正常

问题：服务请求成功率从 99.99% 下降到 99%

问题特征：

- 服务本身 TP 升高，TP9999 大于 50ms（客户端请求成功率下降的直接原因）
 - 请求缓存 TP 指标升高
 - YGC 次数变多
 - YGC 耗时上升，从几毫秒慢慢上升到最高接近 200ms，直到 FullGC
 - 每次 FullGC 性能以上指标变好，然后慢慢变差（每天凌晨 3 点定时主动 FullGC）
- 开始怀疑是 YGC 出现问题。

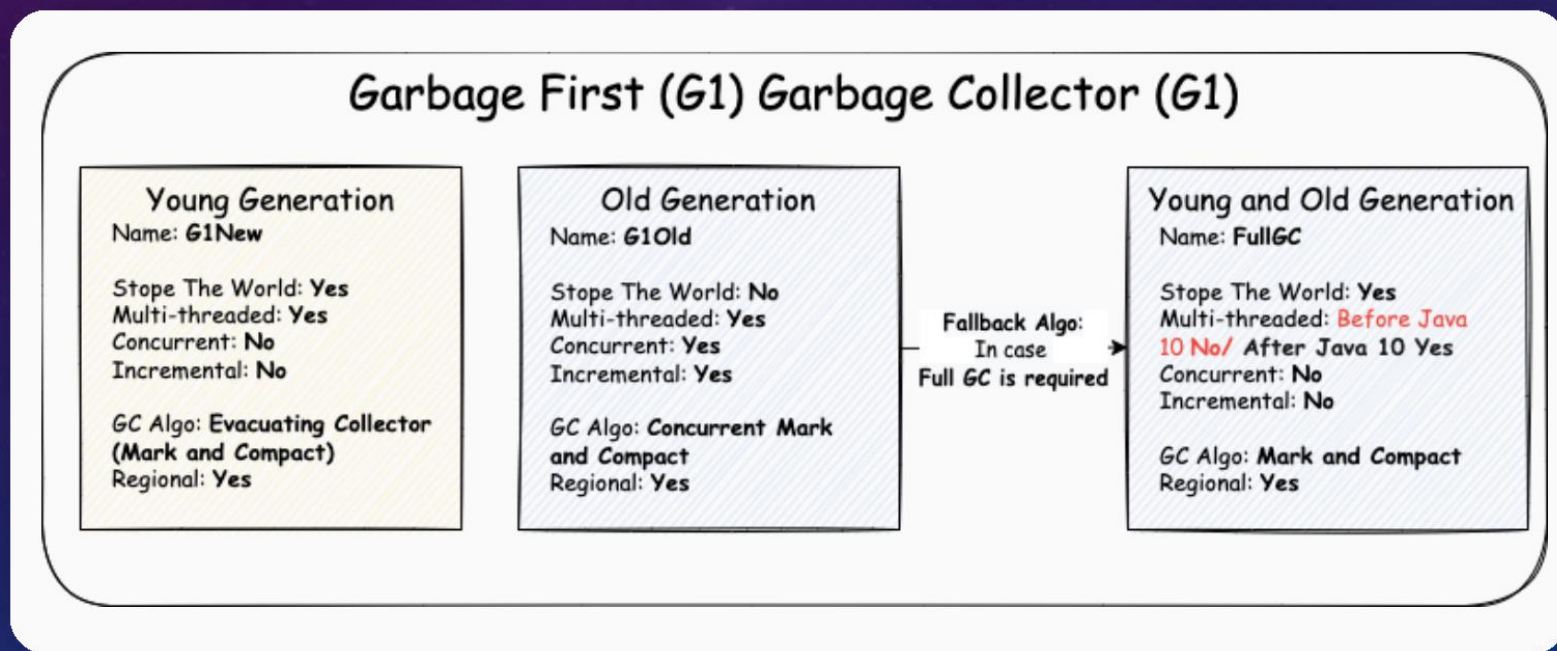
问题：服务请求成功率从 99.99% 下降到 99%

定位到时 YGC 问题后，首先，做止损操作：

- 对机器进行了扩容
- 在午高峰和晚高峰前，各增加一次主动 GC
- 在午高峰和晚高峰期间，紧盯告警，对性能异常的机器进行流量禁用以及手动 FullGC

排查：回顾 G1GC 基本原理


- 分代



排查：回顾 G1GC 基本原理

- 分代
- 分区


G1GC Regions




排查：回顾 G1GC 基本原理

- 分代
- 分区

G1GC Regions




G1GC Young Generation Collection



G1GC Young Generation Collection

排查：回顾 G1GC 基本原理

- 分区
- 分代
- 基于 GC Roots 可达性分析



排查：分析 GC 日志

查看 GC 日志：

- Root Scanning 耗时特别长(正常一般不超过 5ms，现在高达 66ms)

```
74 30.316: [G1Ergonomics (CSet Construction) add young regions to CSet, eden: 468 regions, survivors: 44 regions,
75 30.316: [G1Ergonomics (CSet Construction) finish choosing CSet, eden: 468 regions, survivors: 44 regions, old:
76 , 0.1016196 secs]
77 [Parallel Time: 101.2 ms, GC Workers: 8]
78 [GC Worker Start (ms): Min: 30315.8, Avg: 30315.9, Max: 30315.9, Diff: 0.1]
79 [Ext Root Scanning (ms): Min: 59.8, Avg: 66.5, Max: 68.9, Diff: 9.0, Sum: 532.3]
80 [Update RS (ms): Min: 9.8, Avg: 12.7, Max: 17.0, Diff: 7.3, Sum: 101.6]
81 [Processed Buffers: Min: 1, Avg: 3.6, Max: 15, Diff: 14, Sum: 29]
82 [Scan RS (ms): Min: 0.0, Avg: 0.0, Max: 0.1, Diff: 0.1, Sum: 0.2]
83 [Code Root Scanning (ms): Min: 0.0, Avg: 0.0, Max: 0.0, Diff: 0.0, Sum: 0.0]
84 [Object Copy (ms): Min: 15.2, Avg: 21.5, Max: 30.6, Diff: 15.3, Sum: 172.1]
85 [Termination (ms): Min: 0.0, Avg: 0.2, Max: 1.4, Diff: 1.4, Sum: 2.0]
86 [Termination Attempts: Min: 1, Avg: 300.8, Max: 431, Diff: 430, Sum: 2406]
87 [GC Worker Other (ms): Min: 0.0, Avg: 0.0, Max: 0.0, Diff: 0.0, Sum: 0.2]
88 [GC Worker Total (ms): Min: 101.0, Avg: 101.0, Max: 101.1, Diff: 0.1, Sum: 808.3]
89 [GC Worker End (ms): Min: 30416.9, Avg: 30416.9, Max: 30417.0, Diff: 0.0]
90 [Code Root Fixup: 0.0 ms]
91 [Code Root Purge: 0.0 ms]
92 [Clear CT: 0.1 ms]
93 [Other: 0.3 ms]
94 [Choose CSet: 0.0 ms]
95 [Ref Proc: 0.2 ms]
96 [Ref Enq: 0.0 ms]
97 [Redirty Cards: 0.0 ms]
98 [Humongous Register: 0.0 ms]
99 [Humongous Reclaim: 0.0 ms]
100 [Free CSet: 0.1 ms]
101 [Eden: 468.0M(468.0M)->0.0B(468.0M) Survivors: 45056.0K->45056.0K Heap: 955.5M(1024.0M)->503.5M(1024.0M)]
102 [Times: user=0.69 sys=0.01, real=0.10 secs]
```


排查：分析 GC 日志

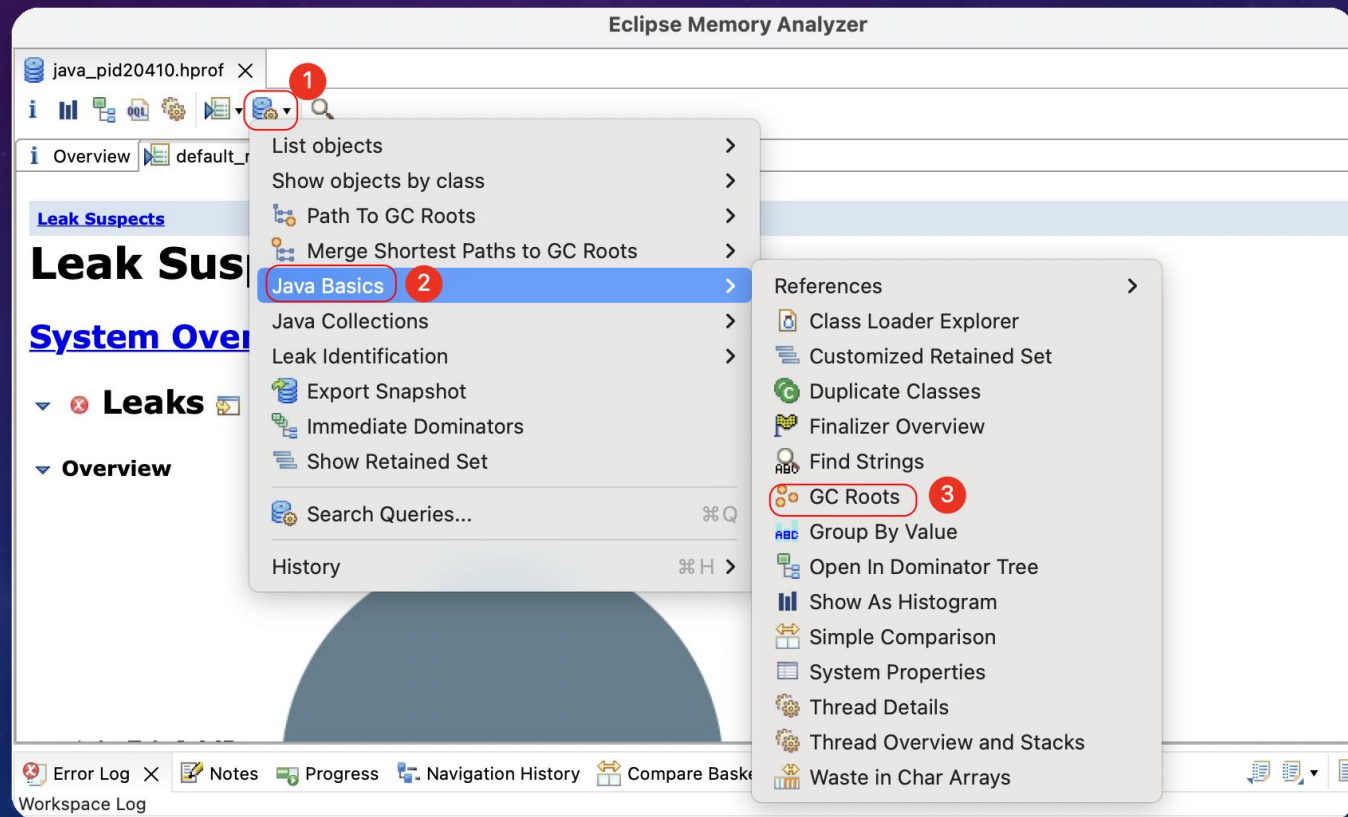
什么是 GC Roots? 网上会告诉你:

- JNI global reference
- JNI local reference
- Local variable on stack
- Monitor
- System class
- Thread
- **Others**

Others 包含哪些取决于 JVM 实现，约等于没说。

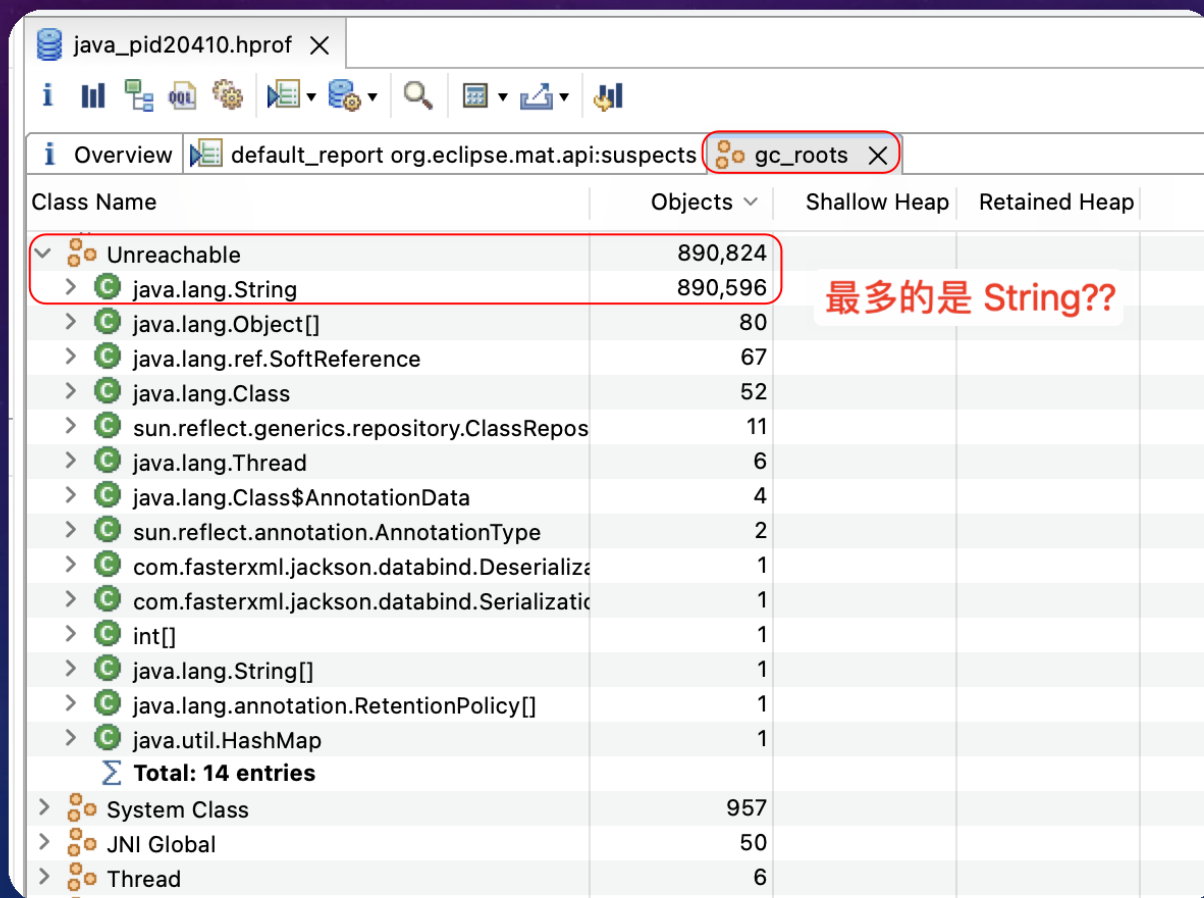
排查：堆内存分析

- Dump 问题现场内存
 - `jmap -dump:format=b,file=dump.hprof <pid>`
 - 这里一定不能加 `live` 参数，否则会先进行 FullGC，现场就丢失了
- 使用 Eclipse Memory Analyzer 分析



排查：堆内存分析

- GC roots 最多的竟然是 String?



java_pid20410.hprof X

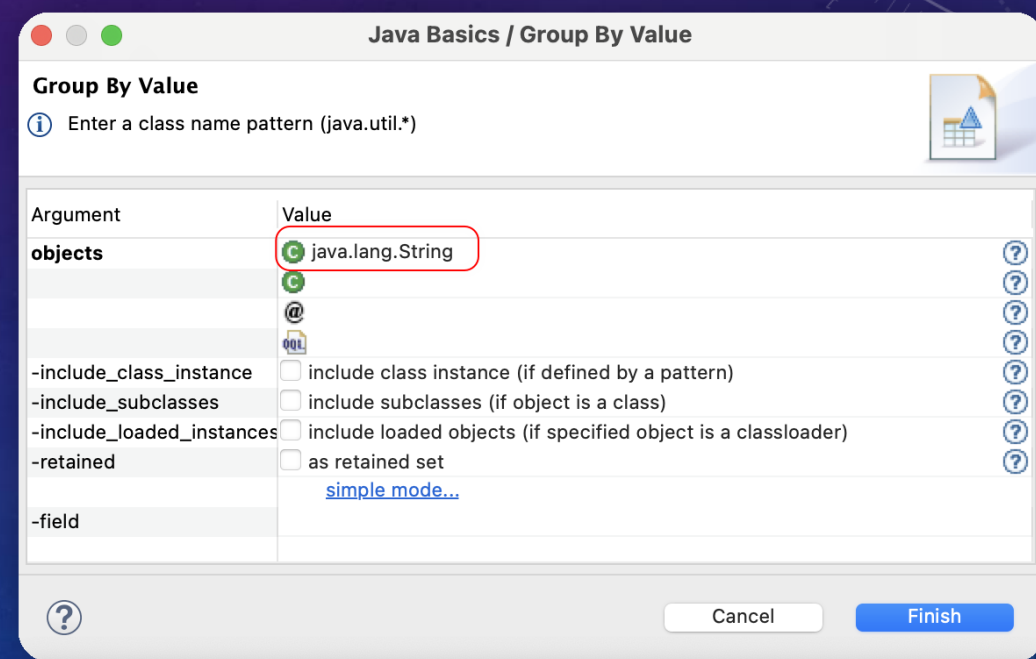
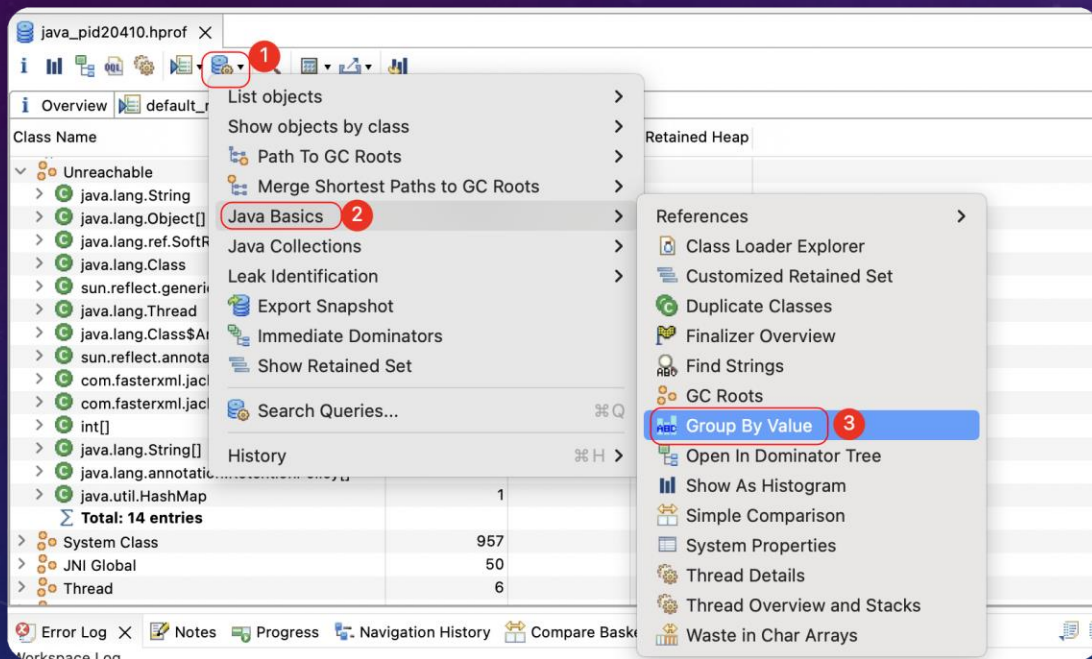
Overview default_report org.eclipse.mat.api:suspects gc_roots X

Class Name	Objects	Shallow Heap	Retained Heap
Unreachable	890,824		
> java.lang.String	890,596		
> java.lang.Object[]	80		
> java.lang.ref.SoftReference	67		
> java.lang.Class	52		
> sun.reflect.generics.repository.ClassRepos	11		
> java.lang.Thread	6		
> java.lang.Class\$AnnotationData	4		
> sun.reflect.annotation.AnnotationType	2		
> com.fasterxml.jackson.databind.Deserializa	1		
> com.fasterxml.jackson.databind.Serializatio	1		
> int[]	1		
> java.lang.String[]	1		
> java.lang.annotation.RetentionPolicy[]	1		
> java.util.HashMap	1		
Σ Total: 14 entries			
> System Class	957		
> JNI Global	50		
> Thread	6		

最多的是 String??

排查：堆内存分析

- 笨办法，使用 MAT 的 Group By Value 功能，看看到底是哪些 String：




排查：堆内存分析

- 发现了上百万个纯数字组成的字符串
 - 每个字符串只出现一次
 - 看起来很像是 userId
- 这些字符串哪里来的呢？
- 为什么字符串会是 GC Roots？

```
tmp > dump > ≡ group_by_string.txt
```

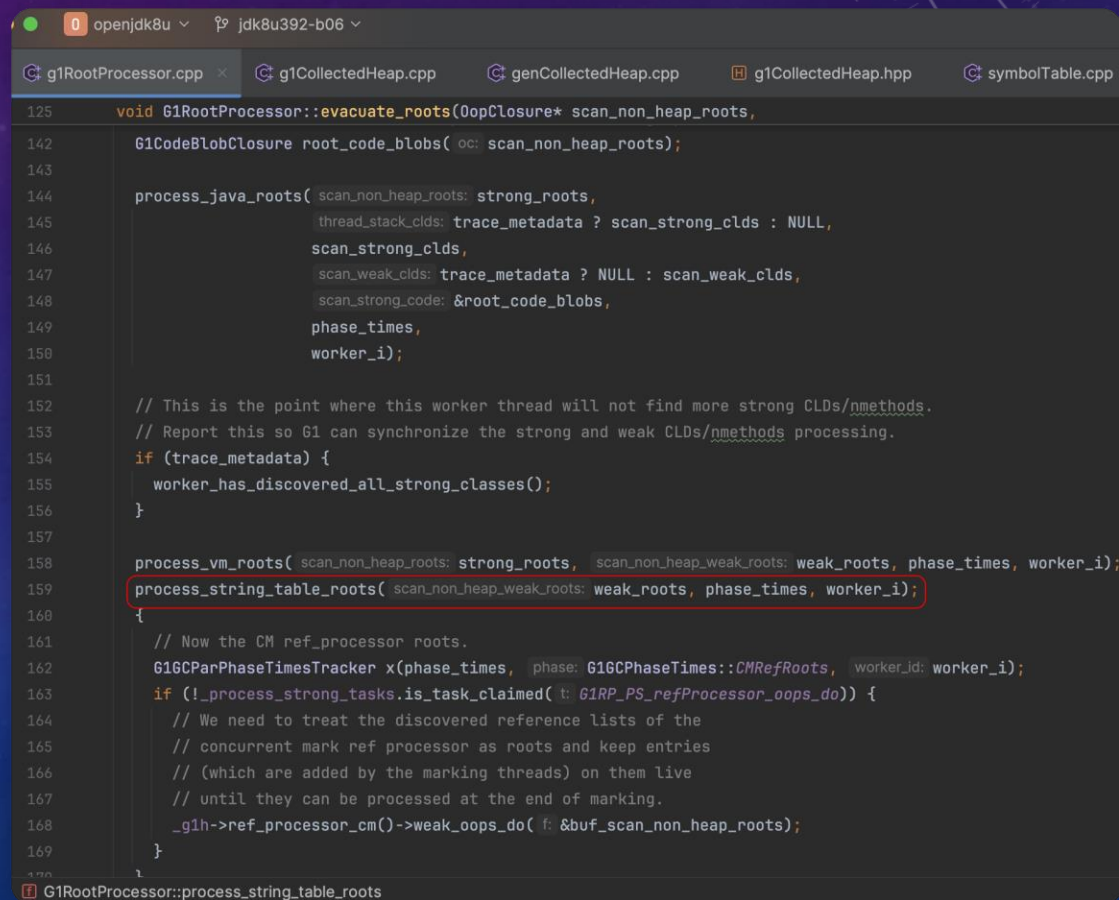
```
570621 6768316651
570622 6768319829
570623 6768329866
570624 6768334868
570625 6768349153
570626 6768358786
570627 6768363519
```

String Value	Objects	Shallow Heap ▾	Avg. Retained Size	Retained Heap
 ^6768**	<Numeric>	<Numeric>	<Numeric>	<Numeric>
6761121963	1	24 B	64 B	64 B
6761122523	1	24 B	64 B	64 B
6761128982	1	34 B	64 B	64 B
6761133729	1	34 B	64 B	64 B
6761139982	1	24 B	64 B	64 B
6761141513	1	24 B	64 B	64 B
6761157718	1	24 B	64 B	64 B
6761158421	1	24 B	64 B	64 B
6761164626	1	24 B	64 B	64 B
6761167155	1	24 B	64 B	64 B
6761167224	1	24 B	64 B	64 B
6761168158	1	24 B	64 B	64 B
6761173618	1	24 B	64 B	64 B

出现次数都是一次

排查：分析 JDK 源码

- G1RootProcessor::evacuate_roots
 - process_java_roots
 - process_vm_roots
 - process_string_table_roots
 - weak_oops_do
 -
- YGC 时，会扫描 StringTable（为什么？）



```
void G1RootProcessor::evacuate_roots(0opClosure* scan_non_heap_roots,  
G1CodeBlobClosure root_code_blobs(oc: scan_non_heap_roots);  
  
process_java_roots( scan_non_heap_roots: strong_roots,  
thread_stack_clds: trace_metadata ? scan_strong_clds : NULL,  
scan_strong_clds,  
scan_weak_clds: trace_metadata ? NULL : scan_weak_clds,  
scan_strong_code: &root_code_blobs,  
phase_times,  
worker_i);  
  
// This is the point where this worker thread will not find more strong CLDs/nmethods.  
// Report this so G1 can synchronize the strong and weak CLDs/nmethods processing.  
if (trace_metadata) {  
worker_has_discovered_all_strong_classes();  
}  
  
process_vm_roots( scan_non_heap_roots: strong_roots, scan_non_heap_weak_roots: weak_roots, phase_times, worker_i);  
process_string_table_roots( scan_non_heap_weak_roots: weak_roots, phase_times, worker_i);  
{  
// Now the CM ref_processor roots.  
G1GCPArPhaseTimesTracker x(phase_times, phase: G1GCPArPhaseTimes::CMRefRoots, worker_id: worker_i);  
if (!_process_strong_tasks.is_task_claimed( t: G1RP_PS_refProcessor_oops_do)) {  
// We need to treat the discovered reference lists of the  
// concurrent mark ref processor as roots and keep entries  
// (which are added by the marking threads) on them live  
// until they can be processed at the end of marking.  
_g1h->ref_processor_cm()->weak_oops_do( f: &buf_scan_non_heap_roots);  
}  
}
```

排查：分析 JDK 源码

- G1RootProcessor::process_string_table_roots

```
void G1RootProcessor::process_string_table_roots(OopClosure* weak_roots, G1GCPhaseTimes* phase_times,
                                                  uint worker_i) {
    assert(weak_roots != NULL, "Should only be called when all roots are processed");

    G1GCParPhaseTimesTracker x(phase_times, phase: G1GCPhaseTimes::StringTableRoots, worker_id: worker_i);
    // All threads execute the following. A specific chunk of buckets
    // from the StringTable are the individual tasks.
    StringTable::possibly_parallel_oops_do(f: weak_roots);
}
```


排查：分析 JDK 源码

- StringTable::possibly_parallel_oops_do

```
913 → void StringTable::possibly_parallel_oops_do(ObjClosure* f) {  
914     const int limit = the_table()->table_size();  
915  
916     for (;;) {  
917         // Grab next set of buckets to scan  
918         int start_idx = Atomic::add(ClaimChunkSize, &_amp;parallel_claimed_idx) - ClaimChunkSize;  
919         if (start_idx >= limit) {  
920             // End of table  
921             break;  
922         }  
923  
924         int end_idx = MIN2(a: limit, b: start_idx + ClaimChunkSize);  
925         buckets_oops_do(f, start_idx, end_idx);  
926     }  
927 }
```


罪魁祸首：？

疑问：这些字符串是怎么进入 StringTable 的？

String Value	Objects	Shallow Heap ▾	Avg. Retained Size	Retained Heap
 ^6768**	<Numeric>	<Numeric>	<Numeric>	<Numeric>
6761121963	1	24 B	64 B	64 B
6761122523	1	24 B	64 B	64 B
6761128982	1	24 B	64 B	64 B
6761133729	1	24 B	64 B	64 B
6761139982	1	24 B	64 B	64 B
6761141513	1	24 B	64 B	64 B
6761157718	1	24 B	64 B	64 B
6761158421	1	24 B	64 B	64 B
6761164626	1	24 B	64 B	64 B
6761167155	1	24 B	64 B	64 B
6761167224	1	24 B	64 B	64 B
6761168158	1	24 B	64 B	64 B
6761173618	1	24 B	64 B	64 B

出现次数都是一次

罪魁祸首：Jackson

- 我们使用 Jackson 进行反序列处理
- Jackson 会对 JSON 的属性名进行 intern 处理
- 早在 2016 年，就有人提了相关 issue: [#332](#)

// 用户请求参数

```
{  
  "6768349153": {  
    "orderId": "xx",  
    "age": 30,  
    "orderTime": 1721195690135  
  }  
}
```

```
InternCache.java  
11 public final class InternCache  
37  
40 public String intern(String input) {  
41     String result = get(input);  
42     if (result != null) { return result; }  
43  
44     /* 18-Sep-2013, tatu: We used to use LinkedHashMap, which has simple LRU  
45     * method. No such functionality exists with CHM; and let's use simplest  
46     * possible limitation: just clear all contents. This because otherwise  
47     * we are simply likely to keep on clearing same, commonly used entries.  
48     */  
49     if (size() >= MAX_ENTRIES) {  
50         /* Not incorrect wrt well-known double-locking anti-pattern because underlying  
51         * storage gives close enough answer to real one here; and we are  
52         * more concerned with flooding than starvation.  
53         */  
54         synchronized (lock) {  
55             if (size() >= MAX_ENTRIES) {  
56                 clear();  
57             }  
58         }  
59     }  
60     result = input.intern();  
61     put(result, result);  
62     return result;  
63 }  
64 }
```

罪魁祸首：Jackson

```
// 用户请求参数
{
  "6768349153": {
    "orderId": "xx",
    "age": 30,
    "orderTime": 1721195690135
  }
}
```

```
1 public class ModelPredictService {
2
3     /**
4      * 伪代码
5      */
6     public <T> T doExecute(String sceneId, String paramJsonStr) {
7         // 1. 解析 JSON 参数
8         final Map<String, Object> paramMap = JacksonJsonParser.parse(paramJsonStr);
9         // 2. 调用特征查询接口
10        final Map<String, Double> featureMap = FeatureQuery.query(sceneId, paramMap);
11        // 3. 调用模型预测接口
12        final double score = ModelPredict.predict(featureMap);
13        // 4. 调用结果处理接口
14        return ResultProcess.process(score, featureMap, paramMap);
15    }
16
17 }
```

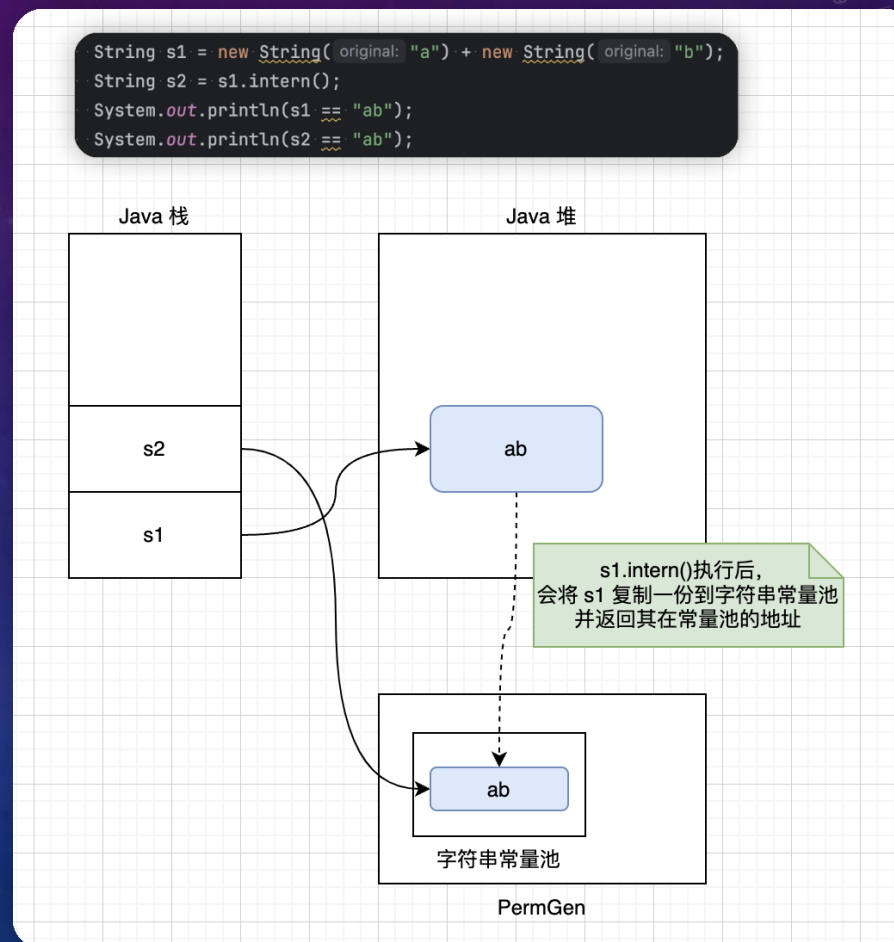
拓展知识：String.intern

- 不『经典』的『面试题』

```
public class StringIntern {  
  
    public static void main(String[] args) {  
        // Java 6 和 Java 7/8 运行结果分别是什么  
        String s1 = new String("a") + new String("b");  
        String s2 = s1.intern();  
        System.out.println(s1 == "ab");  
        System.out.println(s2 == "ab");  
    }  
}
```


拓展知识：String.intern

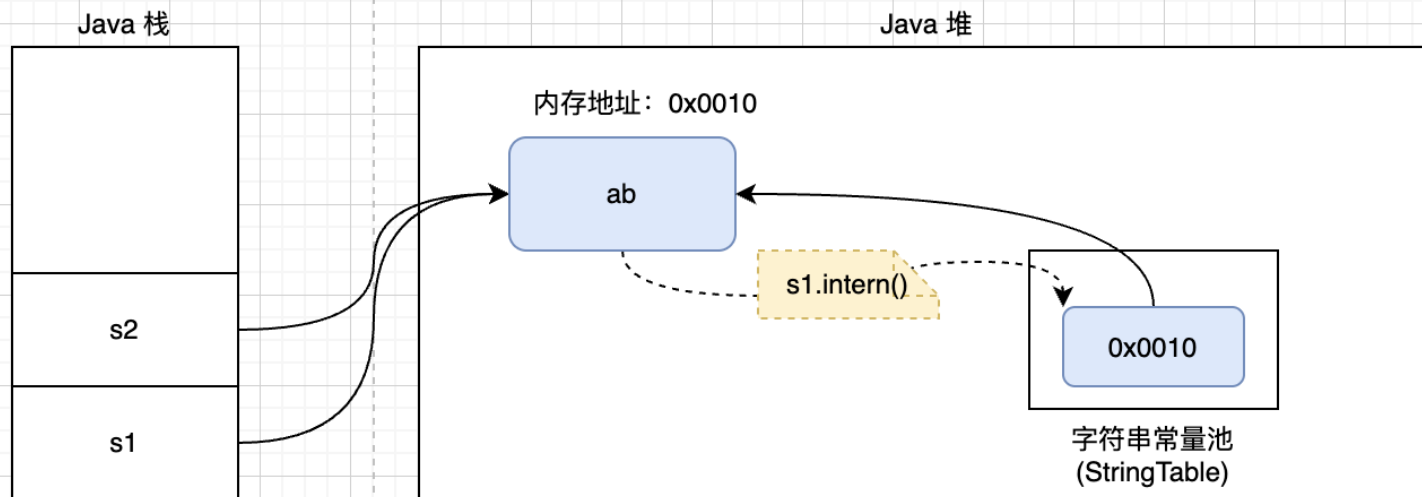
- Java 6
 - 字符串常量池在 PermGen
- 答案
 - `s1 == "ab"` false
 - `s2 == "ab"` true



拓展知识：String.intern

- Java 7/8
 - 字符串常量池在 Java 堆
- 答案
 - `s1 == "ab"` true
 - `s2 == "ab"` true

```
String s1 = new String( original: "a") + new String( original: "b");  
String s2 = s1.intern();  
System.out.println(s1 == "ab");  
System.out.println(s2 == "ab");
```



修复方案：？

思考和分享：怎么修复？

修复方案：禁用 String.intern

修复方案：禁用 Jackson 的 String.intern 功能即可

```
private static final ObjectMapper MAPPER = new ObjectMapper( 1 usage  
    JsonFactory.builder().disable(JsonFactory.Feature.INTER_FIELD_NAMES).build()  
);
```

添加一行代码即可

Feature that determines whether JSON object field names are to be canonicalized using `String.intern` or not: if enabled, all field names will be intern()ed (and caller can count on this being true for all such names); if disabled, no intern()ing is done. There may still be basic canonicalization (that is, same String will be used to represent all identical object property names for a single document).

Note: this setting only has effect if `CANONICALIZE_FIELD_NAMES` is true -- otherwise no canonicalization of any sort is done.

This setting is enabled by default.

```
INTER_FIELD_NAMES( defaultState: true),
```


总结&经验教训

- 遇到问题，首先想止损方案：
 - 如果有变更，尝试回滚
 - 如果扩容可行，扩容
 - 临时增加监控策略，人工运维，优先保证服务可用性
 - ...
- 问题根因：
 - 业务方将类随机字符串用作了 JSON 的属性名，导致 StringTable 膨胀，影响了 YGC
- 问题特征：
 - 量变引起质变，需要足够长的时间，足够多的数据，才能暴露问题
 - 非常隐蔽，伏笔早就埋下，只在调用方改变请求参数时才发生

总结&经验教训

- 排查过程：
 - 尝试录制线上流量进行回放复现问题，失败
 - 分析了 GC 日志，需要熟悉 GC 日志的格式
 - 使用 jmap dump 内存并避免 FullGC
 - 使用 Eclipse Memory Analyzer 对内存dump 进行了分析
 - 查看了 G1 的源码，获得了关键线索

Q&A